Product Task, July 2026

Question 1

About Josh Talks AI

At Josh Talks AI, we work on evaluations and human feedback systems for AI in Indian contexts.

In simple words, this means we care about questions like:

- Is an AI model actually useful for Indian users?
- Does it work well for Indian needs, styles, and contexts?
- Do people prefer one model over another?
- Where does a model perform well, and where does it fail?

A lot of AI systems can look impressive in demos, but real quality becomes clear only when they are tested properly. That is why evaluation is an important part of building better AI products.

Why evaluation matters

AI models can generate outputs, but that alone is not enough.

For example, an image model may generate an image for a prompt, but we still need to ask:

- Is the image relevant?
- Is it visually strong?
- Does it match the user's intent?
- Does it feel right for India?
- Is it better than another model's output?

Good evaluations help teams:

- compare models
- understand what users prefer
- find quality issues
- discover failure cases
- improve products over time

A good evaluation is not only about scoring outputs. It is about creating a system that helps teams make better decisions.

Why this hiring challenge exists

This challenge is designed to test how you think when the problem is not fully defined.

We are intentionally not giving you exact steps, fixed prompts, or a fixed scoring method. We want to see whether you can:

- identify a useful evaluation opportunity
- define it clearly
- run a small version of it
- collect useful evidence
- draw meaningful conclusions

Learn more about evals

Before designing your eval, please look at a few real-world examples to understand how evaluation systems can work in practice:

- Arena AI: https://arena.ai/
- Artificial Analysis: https://artificialanalysis.ai/
- Artificial Analysis Image Arena: https://artificialanalysis.ai/image/arena
- Voice of India: https://voice-of-india.ai.joshtalks.com/
- Josh Talks' Evaluation of Sarvam Bulbul v3

You do not need to copy these. They are only meant to help you understand how real-world evals are structured.

The challenge

Design and run **one useful text-to-image evaluation** for India.

Models to evaluate and how to generate the images

Your evaluation must compare outputs from these 3 required image generation models:

- **OpenAI — GPT Image 1 —** https://developers.openai.com/api/docs/guides/image-generation
- **Google — Gemini 2.5 Flash Image —** https://ai.google.dev/gemini-api/docs/image-generation
- **Google — Gemini 3.1 Flash Image Preview —** https://ai.google.dev/gemini-api/docs/image-generation

You may add other image generation models if you want. If you do, please clearly mention the company name and the exact model name.

Examples:

- **Ideogram — Ideogram 3.0** — https://ideogram.ai/features/3.0

- **Black Forest Labs — FLUX1.1 [pro]** — https://docs.bfl.ai/flux_models/flux_1_1_pro
- **Recraft — Recraft V3** — https://www.recraft.ai/docs/recraft-models/recraft-V3
- **Adobe — Firefly Image 3** — https://blog.adobe.com/en/publish/2024/04/23/adobe-advances-creative-ideation-with-new-firefly-image-3-model

For each model, generate images using the same prompt as far as possible, and keep the setup fair across models.

For every image you submit, save:

- the exact company name
- the exact model name
- the exact prompt
- where you generated it
- any important settings or variations used
- screenshots or proof of generation

If one model requires a slightly different workflow, document that clearly in your submission.

What kind of eval we are looking for

We are **not** looking for a broad benchmark.

We are looking for **one focused and useful evaluation**.

That means:

- choose one **specific category** eg. e-commerce or education
- choose one **specific use case** within that category
- explain why the eval is useful for **India**
- explain why an **AI lab building for India** would care about it

A strong submission should make it easy for us to understand:

- what is being tested
- why it matters
- what "better" means
- why this is worth evaluating

What you need to do

You should:

- define the evaluation
- compare outputs from at least 3 models
- include **human judgment**

- run the evaluation with **8–10 participants**
- show some early findings
- explain how you would scale it further
- also think about the **UI / dashboard / leaderboard** for this evaluation

You can decide the format, judging method, and structure yourself.

Participant requirements

Run the eval with **8–10 participants**.

All participants must be:

- **18+**
- willing to share:
  - name
  - email
  - age
  - consent

You may use this consent statement:

I confirm that I am 18 years or older. I voluntarily participated in this evaluation. I consent to my name, email, and responses/ratings being included in this assignment submission for hiring evaluation purposes.

Deliverables

Please submit all of the following:

1. Main document

This should explain:

- what eval you chose
- what category and use case you selected
- why you chose it
- why it is useful for India
- why it matters for AI labs building for India
- how the evaluation works
- how participants judged the outputs
- what results you found
- how you would scale it further

2. One-page report

Please also submit a **1-page report** that quickly summarizes:

- the eval
- the core setup
- the main findings
- the most important takeaway

This should be written in a way that a busy reviewer can understand quickly.

3. Supporting material

Please include relevant supporting evidence such as:

- prompts
- generated images
- participant ratings
- screenshots
- notes
- consent proof

4. UI / dashboard / leaderboard design

Please also include your thinking on how this eval could look as a product.

This can be simple. For example:

- a basic rating UI
- a dashboard for results
- a leaderboard for model comparison

This does not need to be fully built. A mockup, prototype, Figma, slides, or simple design is fine.

5. Video walkthrough

Please submit a **compulsory video walkthrough** of your work.

The video must:

- be **not more than 2 minutes**
- explain what eval you chose
- explain what you found
- explain why it matters

What we are evaluating

We care about:

- judgment
- originality
- usefulness of the eval

- India-specific thinking
- clarity of reasoning
- quality of execution
- quality of insights
- product thinking

We care much more about **good thinking and strong choices** than polished formatting.

Final reflection

Please end your submission with short answers to these questions:

1. Why did you choose this eval?
2. Why is it useful for India?
3. Why would an AI lab building for India care about it?
4. What did you learn from running the sample?
5. What would you improve with more time?

Question 2

In Question 1, you evaluated AI model outputs using human judgment. Question 2 flips the lens: instead of judging what the AI produces, you'll evaluate the quality of *human* work using data signals instead of human raters. It's the same core skill (defining what "good" means and measuring it), applied to a different target.

Apart from evaluations, Josh Talks AI also builds gold-standard AI datasets across modalities speech, image, and text that help train and fine-tune advanced AI models such as Automatic Speech Recognition (ASR) and Vision-Language (VL) systems.

We operate like a **data research lab**, creating high-quality, accurately labeled datasets that improve the performance of global AI models.

Why ASR and High-Quality Datasets Matter

What is ASR (Automatic Speech Recognition)?

ASR models convert **spoken language (audio)** into **written text** they are the backbone of technologies like:

- Voice assistants (e.g., Alexa, Siri, Google Assistant)
- Meeting transcription tools
- Customer-service bots and call analytics
- Accessibility tools for captioning and dictation
- Voice-based search and commands

These models depend on **training data** pairs of audio and text that teach the AI what words sound like across languages, accents, and tones.

Why Gold-Standard Datasets Matter

An ASR model's quality is measured by its **Word Error Rate (WER)** the percentage of words it gets wrong. Lower WER means higher accuracy. Because the model learns purely from examples, every misheard or mislabeled word in the training data pushes WER up. So models need **large, clean, and accurate datasets** where:

- Audio is recorded in good quality.
- Transcriptions match exactly what's spoken.
- Noise and irrelevant data are removed.
- Annotations are reviewed multiple times for correctness.

Our goal at Josh Talks AI is to create **datasets so clean** that models trained on them achieve **WERs as low as possible (98–99% word-level accuracy)**. In short, data quality directly determines model performance.

Background

As shared above, the newest and fastest-growing division of Josh Talks is working to create datasets to train AI models. This data division is building India's largest-scale, high-quality multilingual voice and vision datasets to train and improve AI models. As you may have seen, AI tools like ChatGPT and others are becoming a big part of our daily lives. Behind the scenes, the quality of the data used to train these models plays a huge role in how well they work — which is why high-quality training data has become more important than ever!

For example, our data engine collects millions of hours of Indian speech, covering 25+ languages, diverse dialects, socio-economic backgrounds, emotional tones, and ambient environments. All data is subjected to **human-in-the-loop quality checks**, ensuring annotation error remains below 2% and delivering enterprise-grade accuracy.

These human review tasks are executed on our in-house platform, Josh Jobs, where we offer freelance opportunities to individuals across India.

The two most common roles include:

- **Recording Tasks:** Contributors submit voice recordings in their native languages, helping us capture 10,000+ hours of audio daily across multiple languages.
- **Transcription & Quality Check Tasks:** Contributors listen to short audio clips in regional languages and transcribe them into accurate text. We currently have over 20,000 daily active users engaged in transcription.

This task is a sample problem statement from a transcription and quality check task.

Here's how transcription works on our platform:

- **Audio Processing:** We have audio chunks (collected from recording) that need accurate text transcriptions.
- **AI First Pass:** We use Whisper AI to create an initial transcription of each audio chunk.

- **Human Review:** Transcribers log into our platform, listen to the audio, and see Whisper's text.
- **Quality Improvement:** They edit Whisper's transcription to fix any errors and create the final, accurate version.
- **Payment:** Transcribers get paid based on the number of accurate hours of audio they complete.

The problem is that some transcribers are doing poor-quality work — either rushing through tasks, making minimal effort, or not actually listening to the audio properly. This hurts our final product quality and wastes money on bad work.

Your Task

Part I

Identify these low-quality transcribers using data patterns. You have access to data of some users and their completed transcription tasks. Each row contains:

- **user_id:** Which person did the work.
- **recording_url:** The audio file they worked on.
- **whisper_text:** AI's transcription of the audio.
- **user_text:** What the transcriber submitted as their final version.
- **is_edited:** Yes if user_text is different from whisper_text (i.e., the user made edits).
- **duration:** How long the audio clip was (in seconds).
- **time_taken_by_user:** How long the transcriber spent on this task (in seconds) (including listening to the audio)
- **segment_character_per_second:** Non-space characters in user_text divided by time_taken_by_user (in seconds).

1. Analyze this data to identify at least 2 warning signs that someone might be doing poor-quality work.
2. For each pattern you identify, explain:

- **What it tells us:** Why does this behavior suggest poor quality?
- **How to measure it:** What would you calculate from the data to spot this pattern?
- **Red flag thresholds:** At what point should we be concerned?
- **False alarms:** When might good transcribers show this pattern innocently?

Part II

Based on your analysis in Part I, design a practical system to automatically spot problematic transcribers on the platform itself, before they do too much damage.

Create clear, actionable recommendations that you could give to our engineering team. Focus on transcription accuracy and transcriber experience, not the technical coding. Keep in mind that the logic/metric you share will be used to block transcribers' accounts, so make sure you are 100% sure about the logic. Your recommendations should look something like this: if a transcriber exhibits X, Y, and/or Z behavior, block the user.