

Final Reflection

1. Why did I choose this eval?

- Public benchmarks (Arena AI, Artificial Analysis) already test general image quality and aesthetics well.
- They do not test a sharper, commercially meaningful question: would a real Indian seller actually post this image on their own store.
- Chose e-commerce catalog photography for solo apparel/handicraft sellers as a narrow, real-world use case rather than a broad benchmark.

2. Why is it useful for India?

- Crores of small sellers depend on Meesho, Instagram, and WhatsApp Business, shooting product photos on their phones with no studio, lighting, or design budget.
- A model that turns a rough photo and description into a sellable product shot is a direct, practical unlock for seller income.
- This use case surfaces India-specific failure modes that generic benchmarks miss: incorrect saree drape, unnatural skin tones, Western-coded styling, garbled text on packaging.

3. Why would an AI lab building for India care about it?

- This is a genuine product wedge, not a demo — the winning model is directly bundle-able into seller-facing tools (Meesho, WhatsApp Business, a Google seller product).
- It tests representation and cultural fit — dimensions a general-purpose leaderboard does not capture.
- Performance here has a direct line to seller adoption and revenue, giving the eval a clear business case.

4. What did I learn from running the sample?

- Gemini 3.1 Flash Image was the clear leader on both average score (3.93/5) and win rate (39%), leading on nearly every axis.
- GPT Image 1 and Gemini 2.5 Flash Image scored identically on average (3.83) but diverged sharply on win rate (27% vs. 34%) — average score and forced-choice pick capture different things.
- “Indian context fit” scored highest of all four axes across every model — a reassuring finding, since this was the axis of greatest concern going in.
- GPT Image 1 scored lowest on prompt relevance specifically, worth a closer qualitative look.
- Model access itself was a finding: GPT Image 1's consumer app no longer serves the exact snapshot, and Gemini 2.5 Flash Image is effectively deprecated on free consumer surfaces and paid-only via API — both required workarounds via LMArena.
- Caught and removed one participant's triplicated submission before computing results — a reminder that data hygiene checks are necessary even in a small, manually-run eval.

5. What would I improve with more time?

- Larger annotator pool — route rating tasks through Josh Jobs to reach hundreds of raters instead of 10.
- More categories beyond apparel/handicrafts — electronics, groceries, services.
- Automated model API pulls so new checkpoints are queued for evaluation automatically.
- An admin audit trail logging configuration and rating changes, needed once this feeds real model-selection decisions.
- Deeper qualitative review of the specific prompts driving GPT Image 1's relevance gap.
- Additional models for comparison — Ideogram 3.0, FLUX1.1 [pro], Recraft V3, Firefly Image 3.