

E-commerce Product Image Evaluation

GPT Image 1 · Gemini 2.5 Flash Image · Gemini 3.1 Flash Image

The Eval

A blind, human-judged comparison of all three required text-to-image models on one focused, commercially real task: generating catalog-ready product photos for small Indian apparel and handicraft sellers on Meesho, Instagram, and WhatsApp Business — the largest social-commerce category in India, where sellers typically have no studio, no lighting, and no design budget.

Core Setup

- **Models:** **GPT Image 1, Gemini 2.5 Flash Image, Gemini 3.1 Flash Image** — all three generated via LMArena's Direct Chat mode, which pins the exact named model version rather than relying on a consumer app's default, ensuring the model actually tested matched the model named in the brief.
- **Prompts:** 7, spanning sarees, jewellery, festive wear, flatlay products, and home décor — avoiding overfit to a single product type.
- **Judging:** 10 consenting participants (18+) rated anonymized, randomized A/B/C outputs on 4 axes — prompt relevance, commercial usability, Indian-context fit, visual quality — plus one forced pick per prompt: “which would you actually post on your store?”
- **Scale:** 10 participants × 7 prompts × 3 options = 210 ratings, 70 forced picks. Data was checked for anomalies before analysis (one duplicate submission found and removed).
- **Infrastructure:** ratings write to a **live Google Sheet in real time via a Google Apps Script Web App** — the same Results dashboard reads directly from that Sheet on every visit, so the leaderboard is always current, not a manual export.

Main Findings

Model	Avg. score (/5)	Win rate	Picks
Gemini 3.1 Flash Image	3.93	39%	27 / 70
GPT Image 1	3.83	27%	19 / 70
Gemini 2.5 Flash Image	3.83	34%	24 / 70

- Gemini 3.1 Flash Image led on both metrics that matter — highest average score and highest win rate — and topped 3 of 4 individual axes.
- GPT Image 1 and Gemini 2.5 Flash Image tied on average score (3.83) but diverged on win rate (27% vs. 34%) — a model can be rated similarly yet chosen less often, a gap a single score would hide.
- “Indian context fit” was the highest-scoring axis for every model (3.84–3.99) — none of the three struggled badly with representation or cultural fit on these prompts.

MOST IMPORTANT TAKEAWAY

Gemini 3.1 Flash Image is the strongest of the three for this use case — not just on average rating, but on the metric that matters most commercially: how often a real seller would actually choose to post the result. The gap between “highly rated” and “actually chosen” (seen clearly in the GPT Image 1 vs. Gemini 2.5 split) is itself the key product insight: win rate, not average score, is the number an AI lab should optimize toward for this kind of seller-facing tool.